

---

---

# COL331 Report: Assignment-3

Manav Rao • 2016CS10523  
Vinayak Rastogi • 2016CS10345

---

# Overview

## Date of Submission

April 25th, 2019

## Report Contents

- Standard operations
  - Create Container
  - Join Container
  - Leave Container
  - Destroy Container
- Command Output
  - ls Command
  - ps Command
  - Copy-On-Write Mechanism
  - Scheduling Mechanism

## Objective:

Implementing container related services in XV6 toy Operating System to understand Virtualization in operating systems.

```

int
sys_create_container(void)
{
    int cid;
    (void)argint(0, &cid);
    return call_create_container(cid);
}

int
sys_destroy_container(void)
{
    int cid;
    (void)argint(0, &cid);
    return call_destroy_container(cid);
}

int
sys_join_container(void)
{
    int cid;
    (void)argint(0, &cid);
    return call_join_container(cid);
}

int
sys_leave_container(void)
{
    return call_leave_container();
}

```

## STANDARD OPERATIONS AND BASIC FUNCTIONALITIES:

- 1) **Create Container :-**
  - a) First inactive/unused slot chosen for initialization
  - b) Changes status to active and return cid of cont.
- 2) **Join Container :-**
  - a) Process is assigned the first empty slot in the cont. w/c it requests to join
  - b) If required a new cont. Can also be created with 0 initial processes
- 3) **Leave Container :-**
  - a) Find cid of the container to w/c process belongs
  - b) Number of Process decreased by 1 after step (a)
  - c) If it were last process, set state of cont. to 'READY'
- 4) **Destroy Container :-**
  - a) Kill all processes of the container
  - b) Container slot state set to UNUSED


\*Container Table -> A separate Table like Data-Structure to keep track of all the containers mimicking how OS keeps track of individual processes.

```
int
call_create_container(int cid)
{
    if(num_containers == 0){
        cstart();
    }
    num_containers++;
    struct inode *ip;
    struct container *c;
    task("Activating the containers");
    for(c = ctable.cont; c < &ctable.cont[NCONT]; c++){
        if(c->state == CUNUSED){
            c->state = CUSED;
            c->cid = cid;
            c->num_procs = 0;
            break;
        }
    }

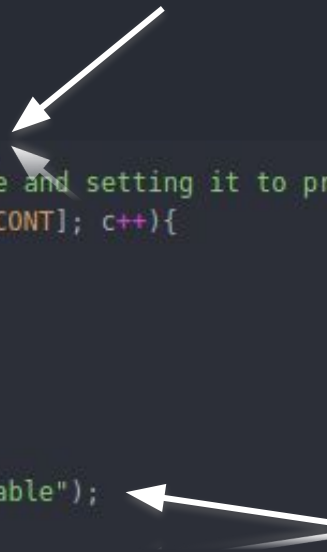
    char path[10];
    path[0] = 'd';
    path[1] = 'i';
    path[2] = 'r';
    path[3] = '\0';

    cprintf("Directory: %s\n", path);
    if ((ip = namei(path)) == 0) {
        return -1;
    }

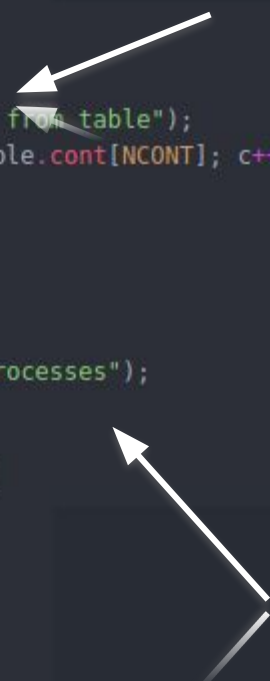
    c->rootdir = ip;
    int i;
    for(i = 0; i < NPROC; i++){
        c->pids[i] = -1;
    }
    return 1;
}
```



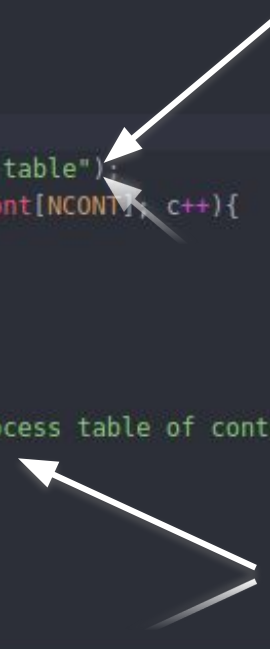
```
int
call_join_container(int cid)
{
    struct proc *p;
    p = myproc();
    struct container *c;
    task("Searching the container from table and setting it to process");
    for(c = ctable.cont; c < &ctable.cont[NCONT]; c++){
        if(c->cid == cid){
            p->cid = cid;
            c->num_procs = c->num_procs + 1;
            break;
        }
    }
    task("Adding process to container pid table");
    int i;
    for(i = 0; i < NPROC; i++){
        if(c->pids[i] == -1){
            c->pids[i] = p->pid;
            break;
        }
    }
    return cid;
}
```



```
int
call_destroy_container(int cid)
{
    if(num_containers == 0){
        return -1;
    }
    num_containers--;
    struct container *c;
    task("Searching the container from table");
    for(c = ctable.cont; c < &ctable.cont[NCONT]; c++){
        if(c->cid == cid){
            c->state = CUNUSED;
            c->cid = -1;
        }
    }
    task("Killing all container processes");
    if(c->num_procs > 0){
        int i;
        for(i = 0; i < NPROC; i++){
            if(c->pids[i] != -1){
                kill(c->pids[i]);
            }
        }
        c->num_procs = 0;
    }
    return 1;
}
```



```
int
call_leave_container(void)
{
    struct proc *p;
    struct container *c;
    p = myproc();
    int cid = p->cid;
    int pid = p->pid;
    p->cid = 0;
    task("Searching the container from table");
    for(c = ctable.cont; c < &ctable.cont[NCONT]; c++){
        if(c->cid == cid){
            break;
        }
    }
    int i;
    task("Removing process from the process table of container");
    for(i = 0; i < NPROC; i++){
        if(c->pids[i] == pid){
            c->pids[i] = -1;
            break;
        }
    }
    return 1;
}
```



---

## COMMAND OUTPUTS OF BASIC SYSTEM CALLS:

### 'ls' command

- Name of the files decided by creator of the files
- The container name is appended to the file name
- When the ls command is called the files which are displayed include:
  - Files created by main (already present files)
  - Files created by the process itself (which are appended by the cid)

### 'ps' command

- It displays the process count as earlier.
- The only difference is that it now displays the process count of a container.
- A check on cid is placed to maintain the same.

```
if(p->state != UNUSED && p->cid == cid){  
    cprintf("pid:%d name:%s\n", p->pid, p->name);
```

---

## COMMAND OUTPUTS OF BASIC SYSTEM CALLS:

### Copy on Write Mechanism

- The private path/name of the file is decided for the container based upon the checks to find if the process is created by the container or not
  - Depending on the Container Id, a separate file with a unique name is derived from the original parent file is created for different containers
  - If the file is already created by another process in the same container then the same file is opened for the current process
  - By this mechanism, this file will not be accessible to other containers
-

```
int
sys_scheduler_log_on(void)
{
    scheduler_log = 1;
    return 1;
}

int
sys_scheduler_log_off(void)
{
    scheduler_log = 0;
    return 1;
}
```

## SCHEDULING MECHANISM:

- A simple kernel end scheduler is maintained.
- A toggle variable is kept to maintain the same.
- In the main scheduling task, the cid is cross checked corresponding to the scheduler\_toggle.